
Trust Region Policy Optimization

TRPO

置信域策略优化

第二部分

东川路第一可爱猫猫虫



主要内容

- $\nabla_{\theta} L_{\theta_{old}}(\theta_{old})$ 的实现
- 用KKT条件求解不等式约束的优化问题
- Hessian Vector Product
- 共轭梯度法求 $Hx = g$
- 回溯线搜索

感谢
蕾西亚万岁啊我透
的续费包月充电

Policy Net

$$\max \ g^T(\theta - \theta_{old})$$

$$s.t. \frac{1}{2}(\theta - \theta_{old})^T H(\theta - \theta_{old}) \leq \delta$$

- g 是 $\nabla_{\theta} L_{\theta_{old}}(\theta_{old})$

$$L_{\theta_{old}}(\theta) = \mathbb{E}_{s \sim \rho_{\theta_{old}}, a \sim \pi_{\theta_{old}}} \left[\frac{\pi_{\theta}(a|s)}{\pi_{\theta_{old}}(a|s)} \cdot A_{\pi_{\theta_{old}}}(s, a) \right]$$

```
log_probs = policy_net.get_log_prob(states, actions)
```

$$\exp(\log \pi_{\theta} - \log \pi_{\theta_{old}}) = \frac{\pi_{\theta}(a|s)}{\pi_{\theta_{old}}(a|s)}$$

```
ratio = torch.exp(log_probs - old_log_probs)
```

g^T



$$L_{\theta_{\text{old}}}(\theta) = \mathbb{E}_{s \sim \rho_{\theta_{\text{old}}}, a \sim \pi_{\theta_{\text{old}}}} \left[\frac{\pi_\theta(a|s)}{\pi_{\theta_{\text{old}}}(a|s)} \cdot A_{\pi_{\theta_{\text{old}}}}(s, a) \right]$$

`loss = (ratio * advantages).mean()`

- g 是 $\nabla_\theta L_{\theta_{old}}(\theta_{old})$, 下一步对 $L_{\theta_{old}}(\theta)$ 求导

`loss_grads = autograd.grad(loss, policy_net.parameters())`

- 把梯度张量展平成一维张量, 拼接得到 g^T

`loss_grad = torch.cat([grad.view(-1) for grad in loss_grads]).detach()`

求解不等式约束的优化问题

$$\max g^T(\theta - \theta_{old})$$

$$s.t. \frac{1}{2}(\theta - \theta_{old})^T H(\theta - \theta_{old}) \leq \delta$$

$$\mathcal{L}(\theta, \lambda) = g^T(\theta - \theta_{old}) - \lambda \left(\frac{1}{2}(\theta - \theta_{old})^T H(\theta - \theta_{old}) - \delta \right)$$

- 对 θ 求偏导并置为0

$$\nabla_\theta \mathcal{L}(\theta, \lambda) = g - \lambda H(\theta - \theta_{old}) = 0$$

$$\theta - \theta_{old} = H^{-1} \cdot \frac{g}{\lambda}$$

- 将 $\theta - \theta_{old}$ 的表达式带入约束条件

$$\frac{1}{2} \left(H^{-1} \cdot \frac{g}{\lambda} \right)^T H \left(H^{-1} \cdot \frac{g}{\lambda} \right) \leq \delta$$

先求 $H^{-1}g$

$$\frac{1}{2} \cdot \frac{g^T H^{-1} g}{\lambda^2} \leq \delta \quad \lambda^2 \geq \frac{g^T H^{-1} g}{2\delta} \quad \lambda \geq \sqrt{\frac{g^T H^{-1} g}{2\delta}}$$

$$\theta - \theta_{old} = H^{-1} \cdot \frac{g}{\sqrt{\frac{g^T H^{-1} g}{2\delta}}} = H^{-1}g \cdot \sqrt{\frac{2\delta}{g^T H^{-1} g}}$$

- 思路是先求 $H^{-1}g$
- 即求解 $Hx=g$
- H 是平均KL散度的Hessian矩阵
- 共轭梯度法

Hessian Vector Product

- 无需显式构造H
就能求出Hessian矩阵和向量的乘积
- Hv
Hessian矩阵与向量 v 的乘积
=一阶导数与向量内积对参数的梯度
- $\theta_{n \times 1}$ 为参数 标量函数 $f(\theta)$

$$\text{一阶导 } g(\theta) = \left[\frac{\partial f}{\partial \theta_1}, \frac{\partial f}{\partial \theta_2}, \dots, \frac{\partial f}{\partial \theta_n} \right]^T$$

$$\text{二阶导 } H(\theta) \quad H_{ij} = \frac{\partial^2 f}{\partial \theta_i \partial \theta_j}$$

下面证明： Hessian矩阵与向量 v 的乘积=一阶导数与向量内积对参数的梯度



Hessian矩阵与向量v的乘积
=一阶导数与向量内积对参数的梯度

- $H_{ij} = \frac{\partial^2 f}{\partial \theta_i \partial \theta_j}$ $v_{n \times 1}$

- $[H(\theta)v]_k$
 $= \sum_{i=1}^n H_{ki} v_i$
 $= \sum_{i=1}^n \frac{\partial^2 f}{\partial \theta_k \partial \theta_i} v_i$

- 一阶导数与向量内积

$$v^T g = \sum_{i=1}^n v_i g_i = \sum_{i=1}^n v_i \frac{\partial f}{\partial \theta_i}$$

求梯度

$$[\nabla_\theta(v^T g)]_k = \frac{\partial(v^T g)}{\partial(\theta_k)} = \sum_{i=1}^n v_i \frac{\partial^2 f}{\partial \theta_k \partial \theta_i}$$

Hvp应用于TRPO 求Hv

- $f(\theta)$ 是KL散度的均值

```
kl = policy_net.get_kl(states)
```

```
kl = kl.mean()
```

- 求 $f(\theta)$ 的一阶导数

```
grads = torch.autograd.grad(kl, policy_net.parameters(), create_graph=True)
```

```
flat_grad_kl = torch.cat([grad.view(-1) for grad in grads])
```

- 一阶导数与向量相乘得到内积

```
kl_v = (flat_grad_kl * v).sum()
```

- 内积对参数求梯度

```
grads = torch.autograd.grad(kl_v, policy_net.parameters())
```

```
flat_grad_grad_kl = torch.cat([grad.contiguous().view(-1)
```

```
for grad in grads]).detach()
```

[DeepRL_Algorithms/Algorith
ms/pytorch/TRPO at
master · RITCHIEHuang/Dee
pRL_Algorithms · GitHub](https://github.com/RITCHIEHuang/DeepRL_Algorithms/tree/master/PyTorch/TRPO)

Hvp



```
def Hvp(v):
    kl = policy_net.get_kl(states)
    kl = kl.mean()
    grads = torch.autograd.grad(kl, policy_net.parameters(),
create_graph=True)
    flat_grad_kl = torch.cat([grad.view(-1) for grad in grads])
    kl_v = (flat_grad_kl * v).sum()
    grads = torch.autograd.grad(kl_v, policy_net.parameters())
    flat_grad_grad_kl = torch.cat([grad.contiguous().view(-1) for grad in
grads]).detach()
    return flat_grad_grad_kl + v * damping
```

共轭梯度 Ax=b

- 初始化

```
x = torch.zeros_like(b, device=device)
```

```
r = - b.clone()
```

```
p = b.clone()
```

```
rdotr = r.t() @ r
```

迭代

- 步长 $\alpha_k = \frac{r_k^T r_k}{p_k^T H p_k}$

- 解的更新

第 $k+1$ 步迭代的解 $x_{k+1} = x_k + \alpha_k p_k$

- 残差的更新

$$\begin{aligned} r_1 &= g - Hx_1 = g - H(x_0 + \alpha_0 p_0) = g - Hx_0 - \alpha_0 H p_0 \\ &= r_0 - \alpha_0 H p_0 \end{aligned}$$

- 更新系数 $\beta_k = \frac{r_{k+1}^T r_{k+1}}{r_k^T r_k}$

- 搜索方向的更新 $p_{k+1} = r_{k+1} + \beta_k p_k$

迭代代码

```
for i in range(steps):
    Hvp = Hvp_f(p) # A @ p
    alpha = rdotr / (p.t() @ Hvp) # step length
    x += alpha * p # update x
    r += alpha * Hvp # new residual
    new_rdotr = r.t() @ r
    betta = new_rdotr / rdotr # beta
    p = - r + betta * p
    rdotr = new_rdotr
if rdotr < rdotr_tol: # satisfy the threshold
    break
```



完整的CG共轭梯度函数

```
def conjugate_gradient(Hvp_f, b, steps=10, rdotr_tol=1e-10):
    x = torch.zeros_like(b, device=device) # initialization approximation of x
    r = - b.clone() # Hvp(x) - b : residual
    p = b.clone() # b - Hvp(x) : steepest descent direction
    rdotr = r.t() @ r # r.T @ r
    for i in range(steps):
        Hvp = Hvp_f(p) # A @ p
        alpha = rdotr / (p.t() @ Hvp) # step length
        x += alpha * p # update x
        r += alpha * Hvp # new residual
        new_rdotr = r.t() @ r
        betta = new_rdotr / rdotr # beta
        p = - r + betta * p
        rdotr = new_rdotr
        if rdotr < rdotr_tol: # satisfy the threshold
            break
    return x
```

求参数的完整更新方向

$$\theta - \theta_{old} = H^{-1} \cdot \frac{g}{\sqrt{\frac{g^T H^{-1} g}{2\delta}}} = H^{-1} g \cdot \sqrt{\frac{2\delta}{g^T H^{-1} g}}$$

```
step_dir = conjugate_gradient(Hvp, loss_grad)
```

```
shs = Hvp(step_dir).t() @ step_dir
```

```
lm = torch.sqrt(2 * max_kl / shs)
```

- 参数的完整更新方向

```
step = lm * step_dir
```



回溯线搜索

```
def line_search(model, f, x, step_dir, expected_improve, max_backtracks=10,
accept_ratio=0.1):
    f_val = f(False).item()
    for step_coefficient in [.5 ** k for k in range(max_backtracks)]:
        x_new = x + step_coefficient * step_dir
        set_flat_params(model, x_new)
        f_val_new = f(False).item()
        actual_improve = f_val_new - f_val
        improve = expected_improve * step_coefficient
        ratio = actual_improve / improve
        if ratio > accept_ratio:
            return True, x_new
    return False, x
```